

Project 2 - Classification with Dimensionality Reduction and Performance Evaluation - Due 02/27/18

Basic requirement (80)

- Task 1 (10): Preprocess the data set. Denote the original dataset as X and the preprocessed data set as nX
 - Dataset used in this project: pima.tr (200 samples) and pima.te (332 samples). They can be downloaded from the Testing Dataset link on the course website, or <http://www.stats.ox.ac.uk/pub/PRNN/>.
 - Each sample in the data set has 7 dimensions. The classification result is 'no': (no diabetes) or 'yes' (diabetes)
 - Refer to the README file for details on the features used
 - Before you can use the dataset, you need to do some preprocessing
 - Change 'yes' and 'no' to 1 and 0 indicating 'with disease' and 'without disease'.
 - Delete the first row in the data set
 - Normalize the data set to make the features comparable (or with the same scale).
 - Suppose x is a sample vector, m_i is the mean of **each feature i** , σ_i is the standard deviation of **each feature i** , then normalization is conducted by $(x - m_i) / \sigma_i$. **Keep in mind that you also need to normalize the samples in the test set. For each sample in the test set, use the same m_i and σ_i you derived from the training set**
- Task 2 (15): Transform the preprocessed dataset using principal component analysis (PCA). Denote the transformed data set as pX .
 - Use PCA to derive a new set of basis and choose the major axes with an error rate **not greater than 0.10**.
 - Represent the data using this new set of basis for a reduced dimension
- Task 3 (15): Using Fisher's linear discriminant (FLD) method to derive the projection direction that best separates the projected data, and generate the projected data. Denote it as fX .
- Task 4 (40): Classification
 - Task 4.1: Use nX . Classify the test set using discriminant functions (Cases I, II, and III) ~~as well as kNN~~.
 - ~~Draw a performance curve with accuracy vs. k values where prior probability is calculated based on the training set.~~
 - Compare the performance of all three ~~four~~ classifiers using prior probability determined by the training set, for fair comparison. **Provide TP, TN, FP, FN values.**
 - **Vary the prior probability and plot sensitivity and specificity with respect to prior probability for the three classifiers.**
 - **Vary the prior probability and plot precision and recall with respect to prior probability**
 - (+15 Bonus for UG) Draw a ROC curve for each of the four classifiers by varying the prior probability. Use the k that generates the best accuracy in the previous drawing.
 - Task 4.2: Repeat Task 4.1 on pX .
 - In addition to the above, plot sensitivity and specificity curves against different error rate (or different numbers of eigenvectors, from 1 to 7)**
 - Task 4.3: Repeat Task 4.1 on fX .

Report (20)